

# LAB EVALUATION - KUSH GUPTA - 4CO26 - 103103734

## Summary of the research paper:

The "Speech Commands: A Dataset for Limited-Vocabulary Speech Recognition" dataset, published by Pete Warden, provides a large corpus of one-second long utterances of short words which essentially act as keywords. The dataset is designed for building and evaluating keyword-spotting systems. It includes more than 65,000 utterances from thousands of different speakers, focusing on recognizing basic commands like "yes," "no," "stop," and others.

## CNN Based multi-class classification of the voice samples:

Code:

```
#import required libraries
import os
import librosa
import numpy as np
import tensorflow as tf
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from tensorflow.keras import layers, models
```

```
print('Modules imported')
```

```
#load and pre-process the data
def load_audio_files_from_folder(folder_path,
sample_rate=16000):
    audio_files = []
    labels = []
    class_names = sorted(os.listdir(folder_path)) # Assuming
folder names are the class labels
```

```
for class_name in class_names:
    class_folder = os.path.join(folder_path, class_name)
    if os.path.isdir(class_folder):
        for file_name in os.listdir(class_folder):
            if file_name.endswith('.wav'):
                file_path = os.path.join(class_folder,
file_name)
                audio, _ = librosa.load(file_path,
sr=sample_rate)
```

```
        # Ensure the audio is 1 second long
        if len(audio) > sample_rate:
            audio = audio[:sample_rate]
        elif len(audio) < sample_rate:
            padding = sample_rate - len(audio)
            audio = np.pad(audio, (0, padding),
'constant')
```

```
        # Normalize the audio
        audio = audio / np.max(np.abs(audio))
```

```
        audio_files.append(audio)
        labels.append(class_name)
```

```
return np.array(audio_files), np.array(labels)
```

```
train_dir = 'speech_commands_v0.02'
audios, labels = load_audio_files_from_folder(train_dir)
print("Audio data shape:", audios.shape)
print("Labels shape:", labels.shape)
print('Data loaded')
```

```
label_encoder = LabelEncoder()
labels_encoded = label_encoder.fit_transform(labels)
```

```
print("Class names:", label_encoder.classes_)
```

```
print("Labels encoded")
```

```
train_audios, val_audios, train_labels, val_labels =
train_test_split(
    audios, labels_encoded, test_size=0.2, random_state=42,
    stratify=labels_encoded
)
```

```
print(f"Training set size: {train_audios.shape}")
print(f"Validation set size: {val_audios.shape}")
```

```
train_dataset =
tf.data.Dataset.from_tensor_slices((train_audios,
train_labels))
val_dataset = tf.data.Dataset.from_tensor_slices((val_audios,
val_labels))
```

```
BATCH_SIZE = 64
train_dataset =
train_dataset.shuffle(1000).batch(BATCH_SIZE).prefetch(tf.data
a.experimental.AUTOTUNE)
val_dataset =
val_dataset.batch(BATCH_SIZE).prefetch(tf.data.experimental.A
UTOTUNE)
```

```
print("Converted to tensorflow dataset")
```

```
model = models.Sequential([
    layers.Input(shape=(16000,)),
    layers.Reshape((16000, 1)),
```

```
    # First convolutional layer
    layers.Conv1D(filters=32, kernel_size=3,
activation='relu', padding='same'),
    layers.MaxPooling1D(pool_size=4),
```

```
    # Second convolutional layer
    layers.Conv1D(filters=64, kernel_size=3,
activation='relu', padding='same'),
    layers.MaxPooling1D(pool_size=4),
```

```
    # Third convolutional layer
    layers.Conv1D(filters=128, kernel_size=3,
activation='relu', padding='same'),
```

```
layers.MaxPooling1D(pool_size=4),
```

```
# Flattening
```

```
layers.Flatten(),
```

```
layers.Dense(128, activation='relu'),
```

```
layers.Dropout(0.3),
```

```
# Output layer
```

```
layers.Dense(len(label_encoder.classes_),  
activation='softmax')
```

```
])
```

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
model.summary()
```

```
print("Model constructed")
```

```
EPOCHS = 10
```

```
history = model.fit(  
    train_dataset,
```

```
        epochs=EPOCHS,  
        validation_data=val_dataset  
)  
print("Model trained")
```

```
model.save('keyword_recognition_cnn_2.h5')  
print("Model saved successfully!")
```

### Data collection of my own voice samples:

Code:

```
import sounddevice as sd  
import wavio  
import os
```

Function to record audio

```
def record_audio(filename, duration=2, fs=16000):  
    print(f"Recording {filename} for {duration} seconds...")  
    recording = sd.rec(int(duration * fs), samplerate=fs,  
channels=1)  
    sd.wait() # Wait until the recording is finished  
    wavio.write(filename, recording, fs, sampwidth=2)  
    print(f"Saved {filename}")
```

```
# Function to create directories and record for each label

def create_dataset(labels, num_samples_per_label=10,
duration=2, dataset_dir="my_custom_dataset"):

    if not os.path.exists(dataset_dir):

        os.makedirs(dataset_dir)
```

```
    for label in labels:

        label_dir = os.path.join(dataset_dir, label)

        if not os.path.exists(label_dir):

            os.makedirs(label_dir)
```

```
        for i in range(num_samples_per_label):

            filename = os.path.join(label_dir, f"{label}
_sample_{i + 1}.wav")

            record_audio(filename, duration=duration)

            print(f"Recorded and saved {filename}")
```

```
# List of labels/words you want to record

labels =

['_background_noise_', 'backward', 'bed', 'bird', 'cat', 'dog', 'do
wn', 'eight', 'five', 'follow', 'forward', 'four', 'go',
'happy', 'house', 'learn', 'left', 'marvin', 'nine', 'no',
'off', 'on', 'one', 'right', 'seven', 'sheila', 'six', 'stop',
'three', 'tree', 'two', 'up', 'visual', 'wow', 'yes', 'zero']

create_dataset(labels, num_samples_per_label=30) # Record 5
samples per word
```

Fine Tuning the model on my own data:

Code:

```
import os
import numpy as np
import librosa
import tensorflow as tf

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import classification_report,
confusion_matrix, accuracy_score, precision_score,
recall_score, f1_score

def load_audio_files_from_folder(folder_path,
sample_rate=16000):
    audio_files = []
    labels = []
    class_names = sorted(os.listdir(folder_path))

    for class_name in class_names:
        class_folder = os.path.join(folder_path, class_name)
        if os.path.isdir(class_folder):
            for file_name in os.listdir(class_folder):
                if file_name.endswith('.wav'):
                    file_path = os.path.join(class_folder,
file_name)
                    audio, _ = librosa.load(file_path,
sr=sample_rate)

# Ensure audio is 1 second long
```



```
        if len(audio) > sample_rate:
            audio = audio[:sample_rate]
        elif len(audio) < sample_rate:
            padding = sample_rate - len(audio)
            audio = np.pad(audio, (0, padding),
'constant')
```

```
        # Normalize audio
        audio = audio / np.max(np.abs(audio))
```

```
        audio_files.append(audio)
        labels.append(class_name)
```

```
    return np.array(audio_files), np.array(labels)
```

```
dataset_dir = 'my_custom_dataset'
audios, labels = load_audio_files_from_folder(dataset_dir)
```

```
print("Audio data shape:", audios.shape)
print("Labels shape:", labels.shape)
print('Pre processing done')
```

```
from sklearn.preprocessing import LabelEncoder
```

```
label_encoder = LabelEncoder()
```

```
labels_encoded = label_encoder.fit_transform(labels)
```

```
print("Class names:", label_encoder.classes_)  
print('Labels encoded')
```

```
from sklearn.model_selection import train_test_split
```

```
train_audios, val_audios, train_labels, val_labels =  
train_test_split(  
    audios, labels_encoded, test_size=0.2, random_state=42,  
    stratify=labels_encoded  
)
```

```
print(f"Training set size: {train_audios.shape}")  
print(f"Validation set size: {val_audios.shape}")  
print('splitted data')
```

```
import tensorflow as tf  
from tensorflow.keras import layers, models
```

```
pretrained_model =  
tf.keras.models.load_model('keyword_recognition_cnn.h5')
```

```
for layer in pretrained_model.layers[:-2]:  
    layer.trainable = False
```

```
model = models.Sequential(pretrained_model.layers[:-2])
```

```
model.add(layers.Flatten(name='new_flatten'))  
model.add(layers.Dense(128, activation='relu',  
name='new_dense_1'))  
model.add(layers.Dropout(0.3, name='new_dropout'))  
model.add(layers.Dense(len(label_encoder.classes_),  
activation='softmax', name='new_output'))
```

```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
model.summary()
```

```
print('model trained')
```

```
train_dataset =  
tf.data.Dataset.from_tensor_slices((train_audios,  
train_labels))  
val_dataset = tf.data.Dataset.from_tensor_slices((val_audios,  
val_labels))
```

```
BATCH_SIZE = 32
train_dataset =
train_dataset.shuffle(1000).batch(BATCH_SIZE).prefetch(tf.data
a.experimental.AUTOTUNE)
val_dataset =
val_dataset.batch(BATCH_SIZE).prefetch(tf.data.experimental.A
UTOTUNE)
```

```
EPOCHS = 100
```

```
history = model.fit(
    train_dataset,
    epochs=EPOCHS,
    validation_data=val_dataset
)
print('Model trained')
```

```
model.save('finetuned_model.h5')
print("Model fine-tuned and saved successfully!")
```

```
val_loss, val_acc = model.evaluate(val_dataset)
print(f"Validation Loss: {val_loss}")
print(f"Validation Accuracy: {val_acc}")
```

```
# Get predictions from the model
y_true = []
y_pred = []
```

```
for x, y in val_dataset:
    y_true.extend(y.numpy()) # Get true labels
    y_pred.extend(np.argmax(model.predict(x), axis=1)) # Get
predicted labels
```

```
y_true = np.array(y_true)
y_pred = np.array(y_pred)
```

```
# Calculate metrics
```

```
# Accuracy
accuracy = accuracy_score(y_true, y_pred)
print(f"Accuracy: {accuracy}")
```

```
# Precision, Recall, and F1-score
precision = precision_score(y_true, y_pred,
average='weighted')
recall = recall_score(y_true, y_pred, average='weighted')
f1 = f1_score(y_true, y_pred, average='weighted')
```

```
print(f"Precision: {precision}")
print(f"Recall: {recall}")
print(f"F1 Score: {f1}")
```

```
# Classification Report
print("\nClassification Report:")
print(classification_report(y_true, y_pred,
target_names=label_encoder.classes_))
```

```
# Confusion Matrix
print("\nConfusion Matrix:")
conf_matrix = confusion_matrix(y_true, y_pred)
print(conf_matrix)
print('Metrics calculated')
```